



# About My Coworkers

When AI Joined My Personal Development Team

A candid look at how artificial intelligence is transforming the developer experience.  
And why you might be missing out.





# About My Coworkers

## When AI Joined My Personal Development Team

A candid look at how artificial intelligence is transforming the developer experience [and why you might be missing out.](#)

# Thank You Sponsors!



## Titanium Sponsors



## Gold Sponsors



## Other Awesome Organizations



# Nick Nisi

- Developer Experience at [WorkOS](#)
- Organizer at [NebraskaJS](#)
- Podcaster
  - formerly of [JS Party](#), now [Dysfunctional](#)
  - [Software Engineering Daily](#)
- Emcee at [SquiggleConf](#) next month
- AI Engineer?

Find me at [vim.dad](#)

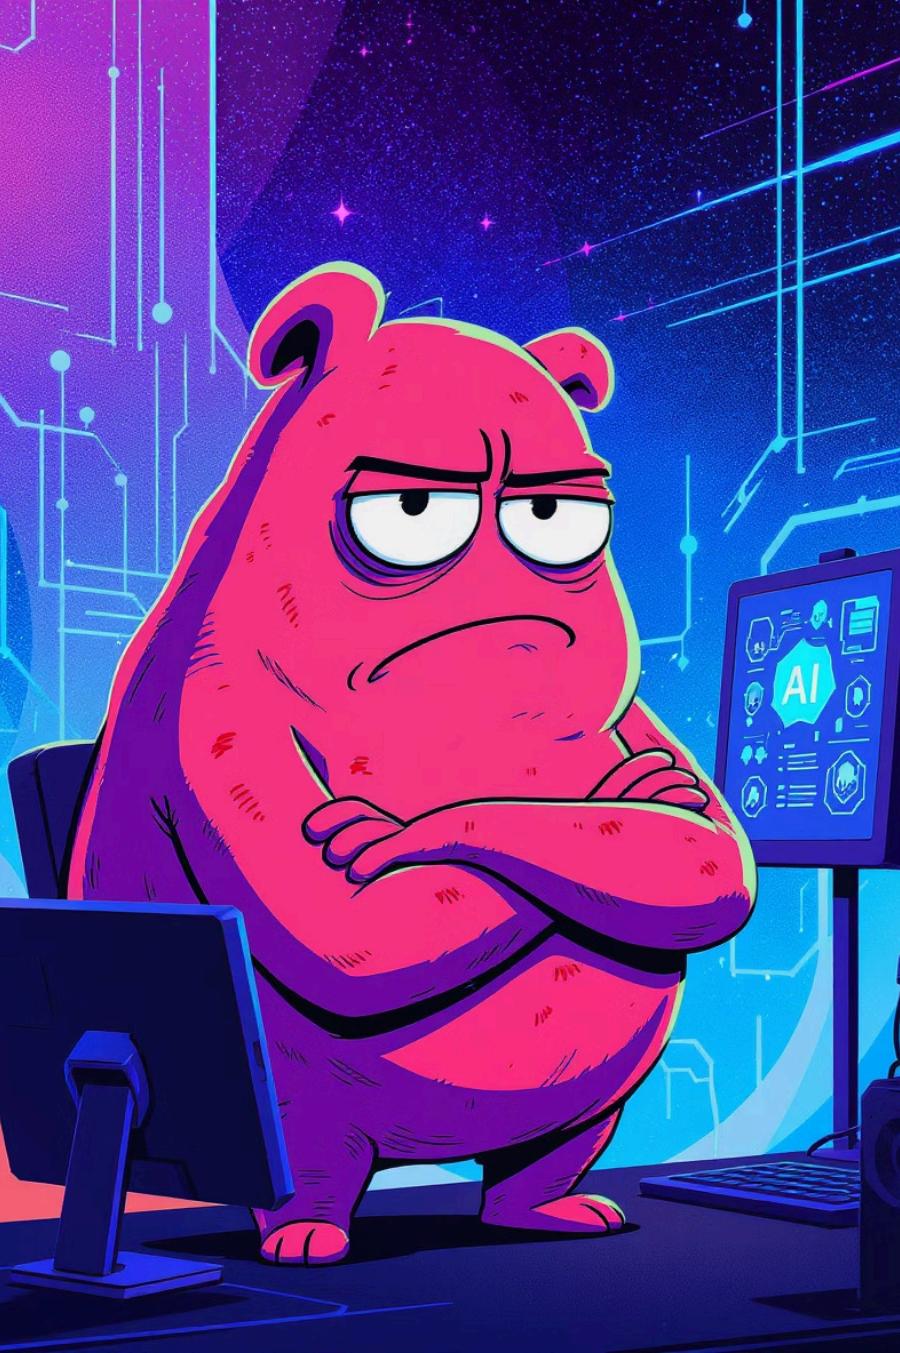




If you're not experimenting  
with AI  
**TODAY**  
You're  
**ALREADY BEHIND.**



**(That was meant to be provocative)**



# I also used to be skeptical

It just wasn't very good.

- Experimenting with Copilot
- Occasionally asking ChatGPT a question
- Whatever Meta had
- Felt like using it was "cheating"



# JS Nation / React Summit

- These amazing developers weren't hiding anything
- They were **collaborating** with AI
- While I was hiding, they were innovating



**That was my wakeup call.**





**BEFORE  
ASKING  
FOR MORE  
HEADCOUNT  
& RESOURCES,  
TEAMS MUST  
DEMONSTRATE  
WHY THEY  
CANNOT GET  
WHAT THEY  
WANT DONE  
USING AI**

**TOBIAS LÜTKE**  
CEO, SHOPIFY



## I'm not the only one...

- AI mandate at Shopify
- Other companies too



The writing is on the wall – **AI adoption isn't optional anymore.**



# Are we overreacting?

Maybe.





"If AI can solve your problems, you don't have hard problems."



What they're really saying is

# "AI didn't work for me"

- Was their problem really that hard? **Maybe.**
- Were they just not good at structuring the problem for AI? **Probably.**
- Does this mean AI is bad and will never be good? **NOPE.**





**This is the worst it'll ever be**

**We still need to learn how to engage it and when not to use  
it.**





# The AI Engineer Mindset

- The tools come and go, but the concepts remain
- AI isn't meant to **replace** you, it's meant to **augment** you.
- Your responsibility definitely hasn't increased. But it has **evolved**.



# AI Likely won't take your job



**Someone using AI might though**

The real competition isn't AI, it's those who embrace and leverage it effectively.

**Ready or not, everything's about to change**

Your daily workflow is about to look completely different. Are you prepared to adapt?

**That spark you felt when you first coded? Will it survive?**

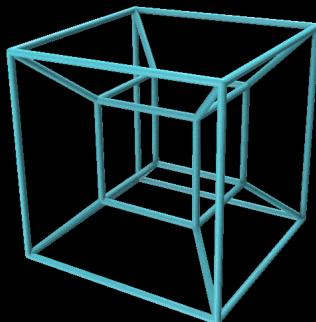
Prepare to rethink what your ideal job entails as AI reshapes daily tasks and responsibilities.

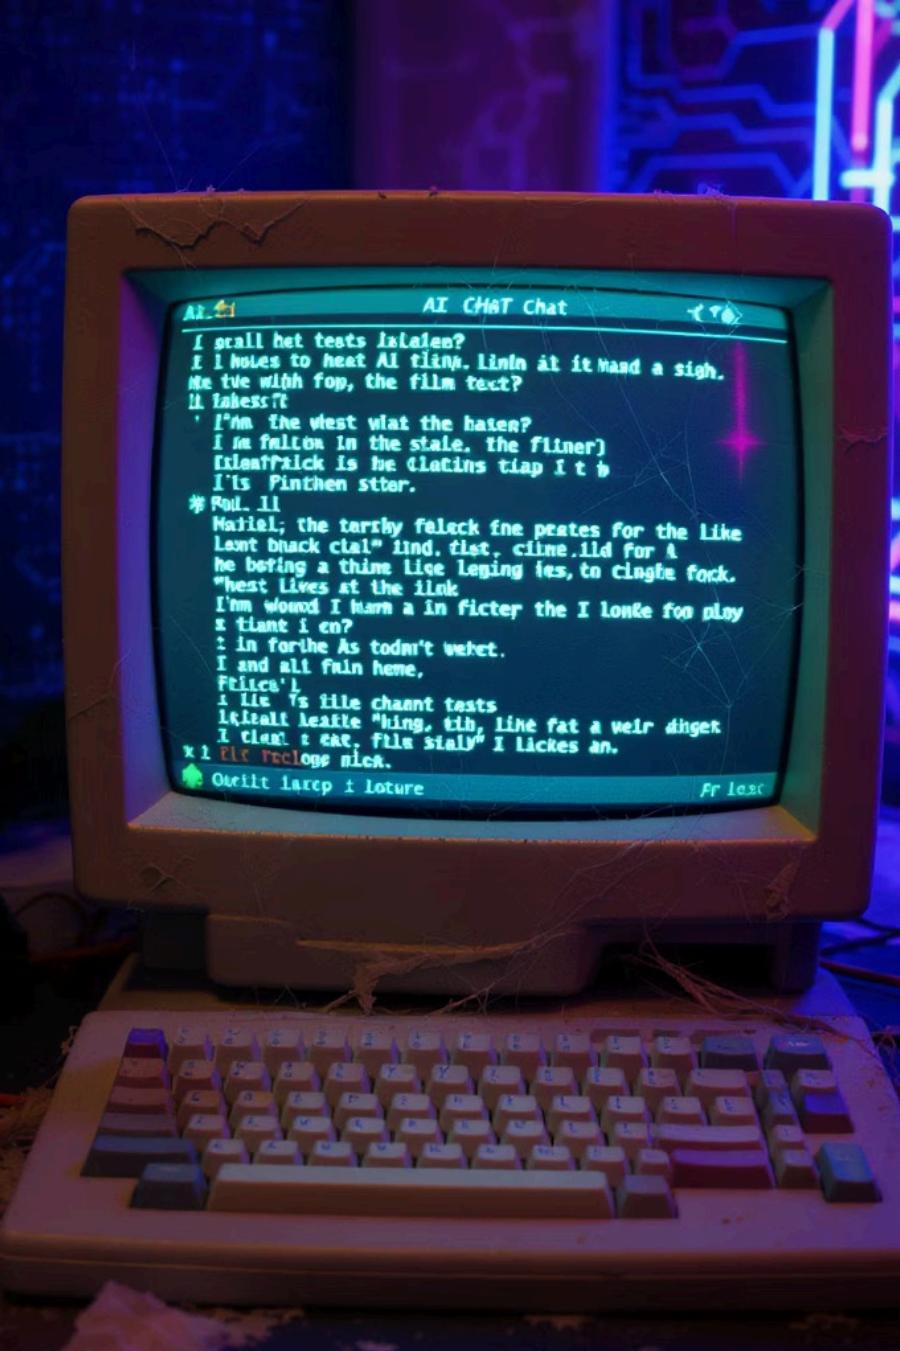




copy / paste

# The dimensions of AI Adoption



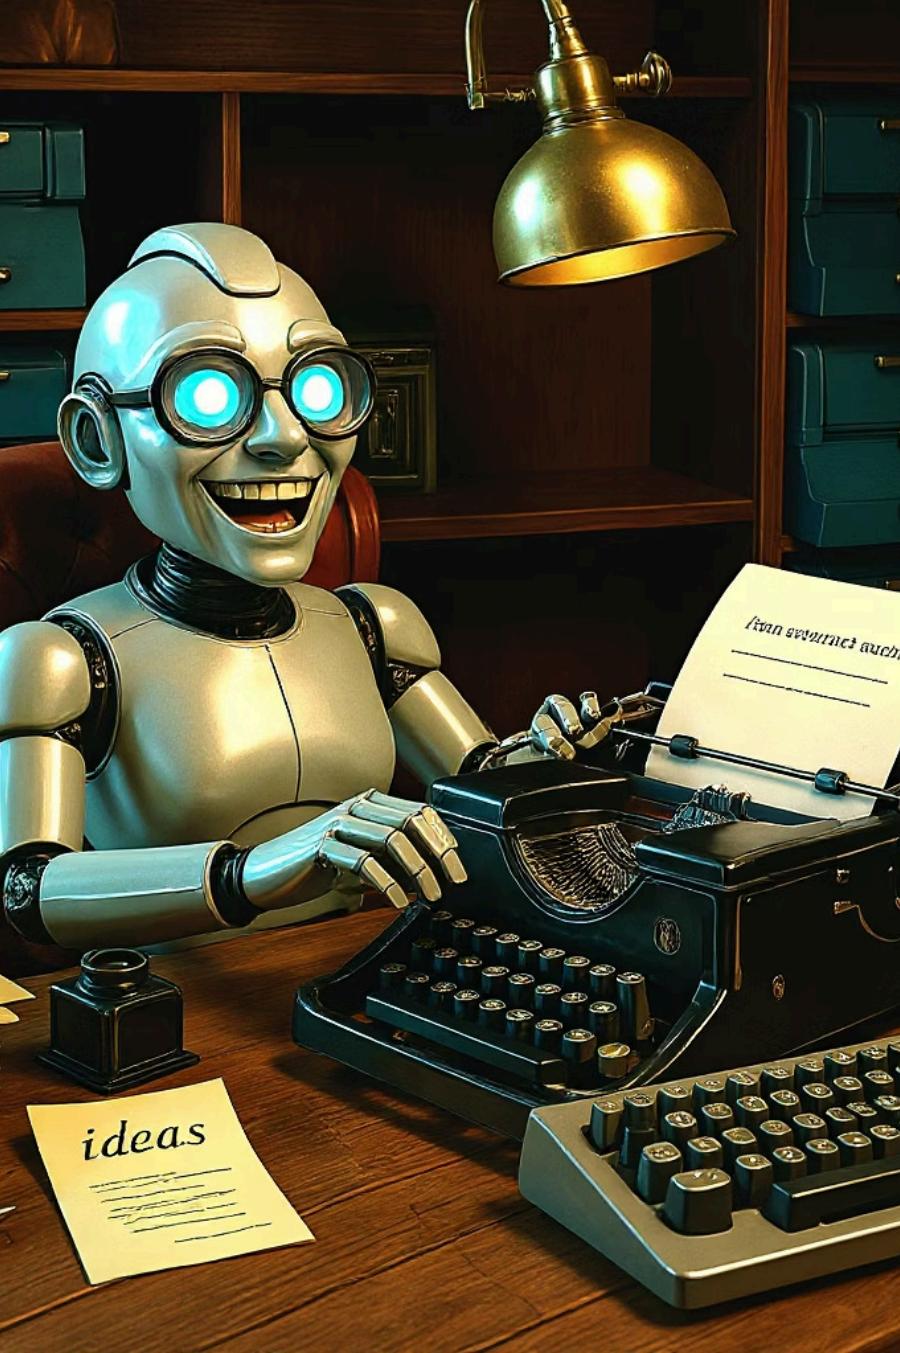


# 1st Dimension: Smarter Autocomplete

- GitHub Copilot suggests the next line
- You hit tab to accept
- Still basically autocomplete, just smarter

This autocomplete is smart - I can kind of chat with it via comments.





## 2nd Dimension: Having Conversations

- Copy code into ChatGPT
- Ask contextual questions about the code

I can chat with this - Maybe I can paste in my file



It's manual.

It's tedious.

But it works.

The "copy-paste-chat" method, while effective, still requires significant human effort and can be repetitive. It's a stepping stone to more integrated AI workflows.



# 3rd Dimension: Contextual Understanding

- Moving beyond just function or single-file context
- AI understands your codebase
- Still lots of copy/paste
- [gitingest.com](https://gitingest.com) - Contextualize and paste an entire codebase

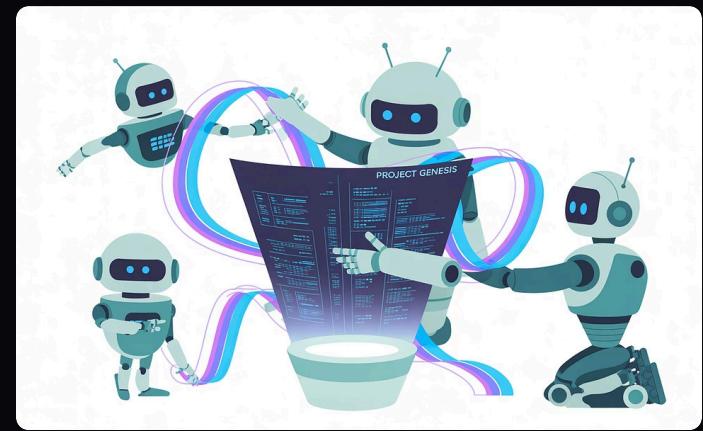
You're the messenger, copying and pasting, back and forth, to manage the conversation.





This is where a lot of people stop.

# The real power: Dimensions 4-7



## 4th Dimension: Pairing with an Agent

AI moves beyond reactive responses, anticipating needs, suggesting architectural improvements, and actively debugging your code before you even ask. It's like having a hyper-intelligent pair programmer.

## 5th Dimension: Agents with Tools

AI takes on larger, end-to-end tasks with minimal human oversight. Think generating entire modules, writing comprehensive test suites for new features, or refactoring vast sections of code on its own.

## 6th Dimension: Agents and other Agents

Specialized AI agents work in concert, communicating and coordinating efforts across different aspects of a project. A dedicated AI for UI, another for backend logic, and one for rigorous testing, all orchestrated seamlessly.

## Nexus



## 7th Dimension: Agent Orchestration

AI becomes a true creative partner, not just automating but generating novel solutions, optimizing algorithms beyond human capacity, and even proposing new product features based on complex data analysis.



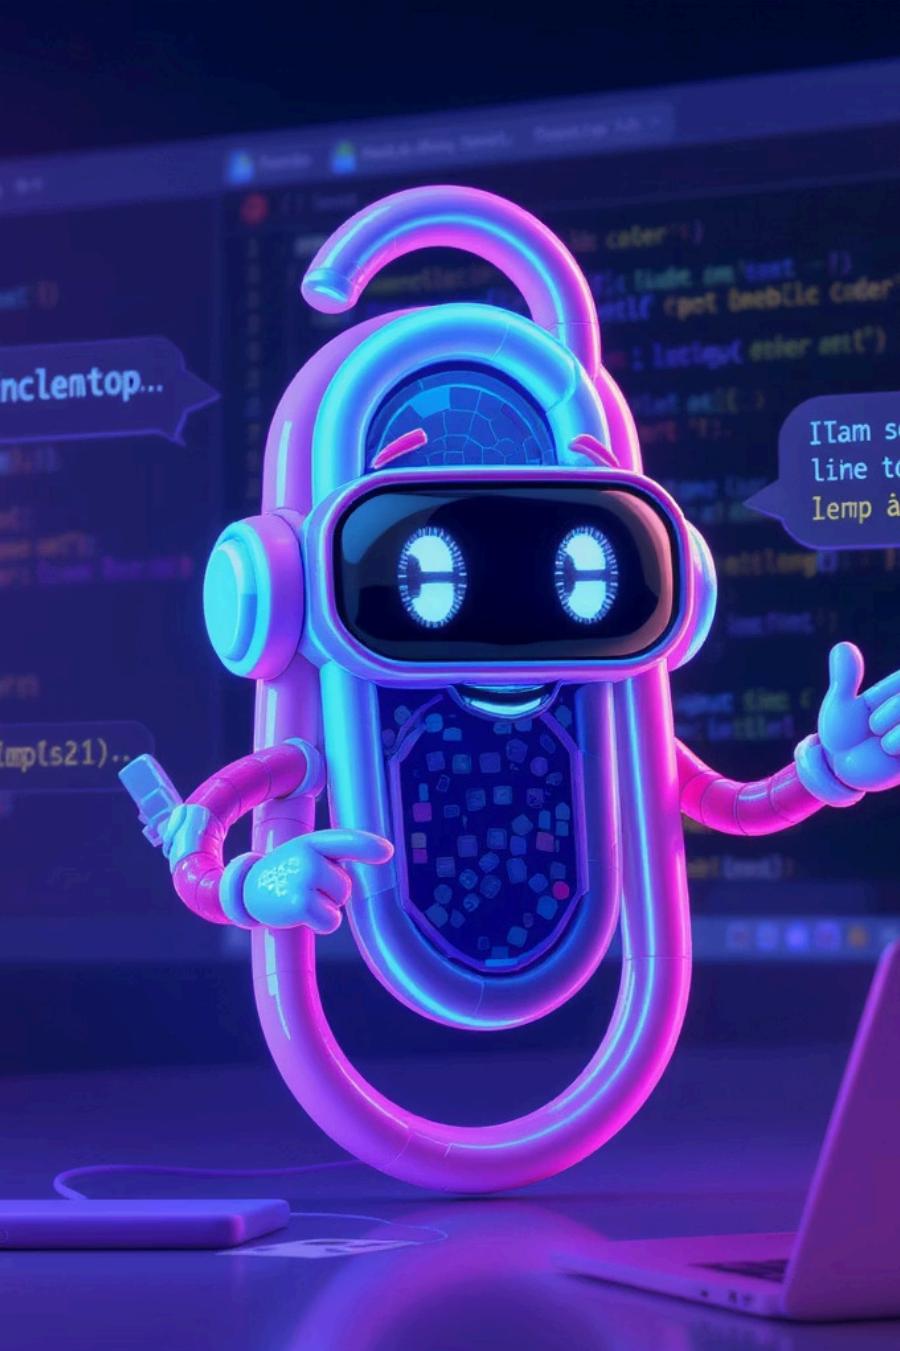


# Dimensions 4-7

## AI Your True Collaborator

Beyond simple assistance, the higher dimensions of AI adoption unlock a powerful partnership where AI actively contributes to problem-solving and innovation.





# 4th Dimension: Pairing with an Agent

Claude Code, Cursor, Amp, etc.

- No more copy/paste
- AI can see your files directly
- AI can edit your files directly

Passive assistance becomes active participation



# 5th Dimension: Agents with Tools

- Agents run tests
- Execute commands
- See the results
- MCP

The AI doesn't tell you what to do - it can run the tools and react to the results.

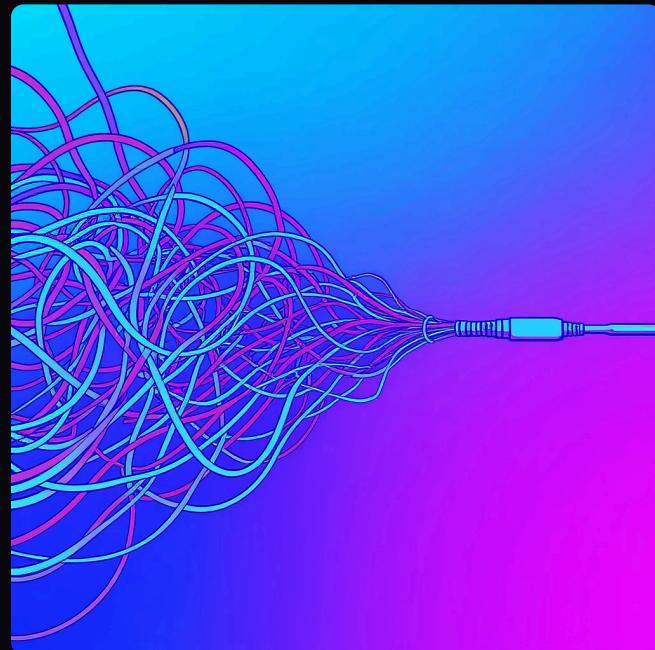




# Model Context Protocol (MCP)

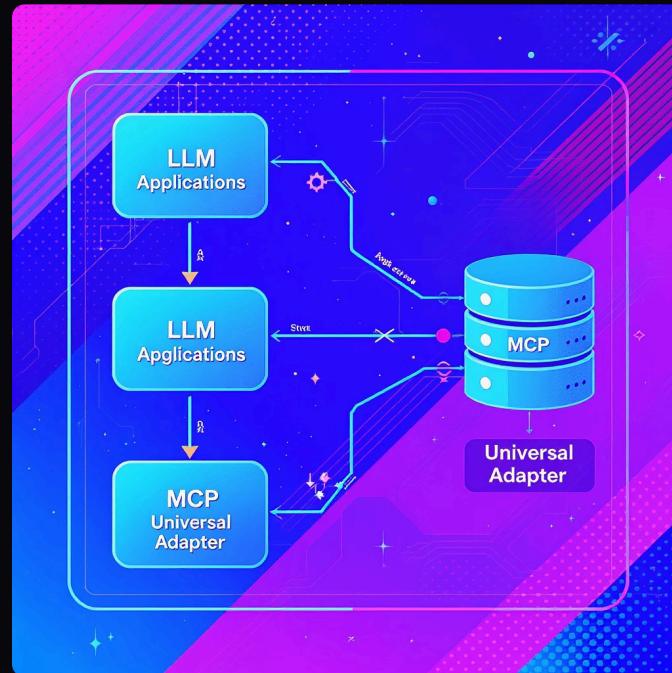
An open standard for seamlessly connecting Large Language Models (LLMs) to various external data sources and tools.

## Why MCP?



Eliminates the " $N \times M$  problem" of redundant integrations by standardizing how AI applications connect to data and tools.

## How it Works



Functions as a client-server architecture, where LLM apps (clients) connect to MCP servers (data/tools), much like a USB-C for AI applications.

## Key Benefits

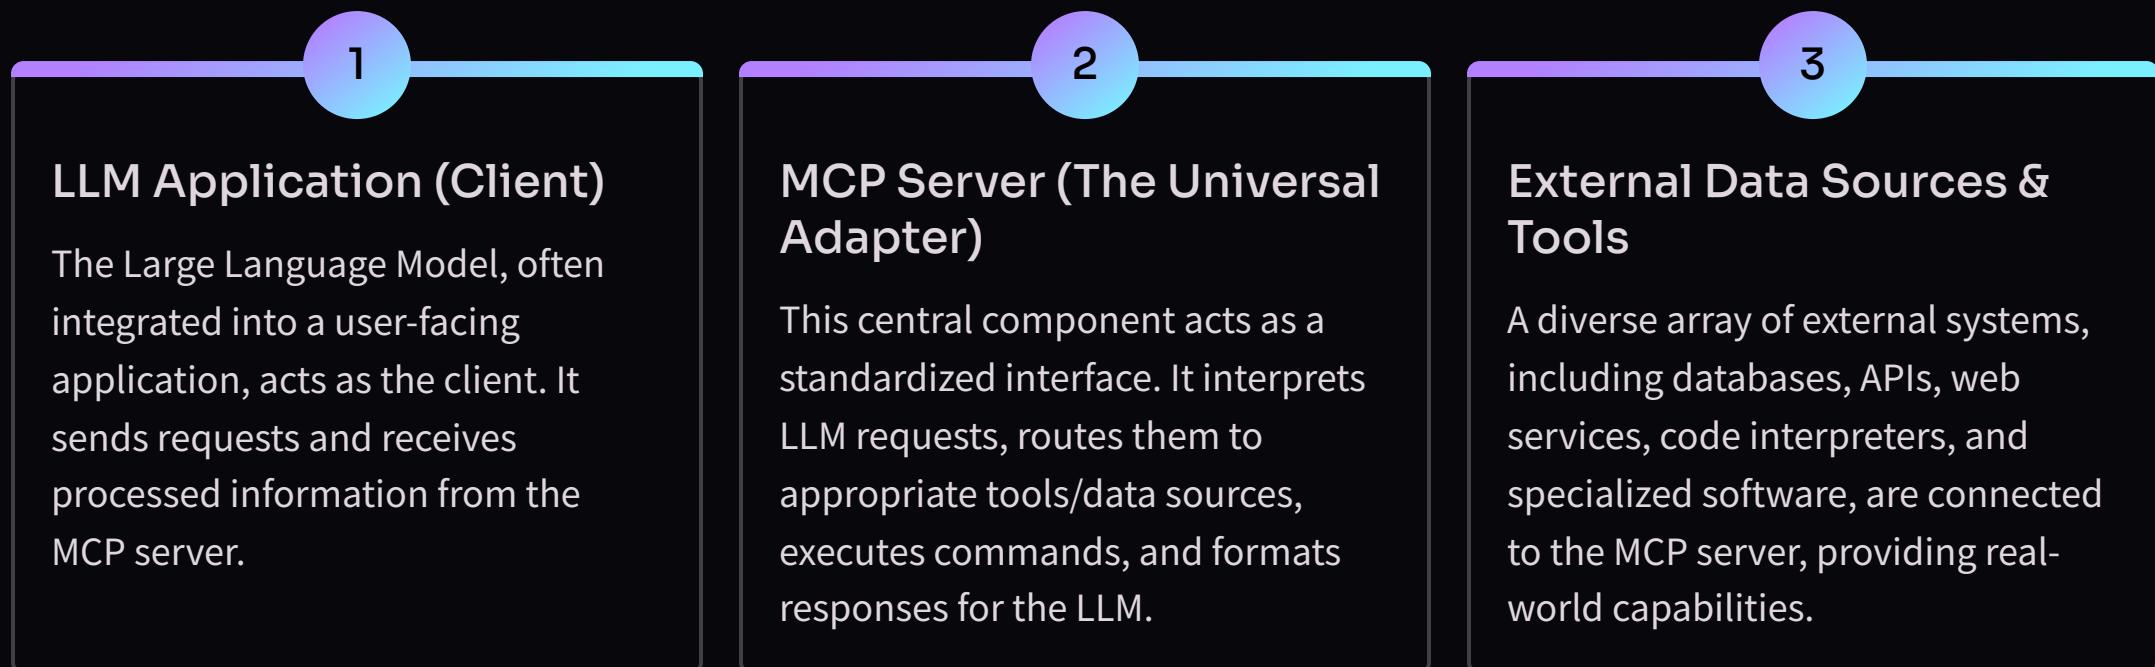


Offers real-time data access, multi-source connectivity, and significantly simplified integration for AI applications.



# The Architecture of MCP: Bridging AI and the World

The Model Context Protocol (MCP) operates as a universal translator, enabling Large Language Models (LLMs) to interact dynamically with the vast ecosystem of existing software tools and data, breaking down traditional barriers and expanding AI capabilities.





# From Passive to Active

Large Language Models (LLMs) are evolving beyond mere information providers, now capable of executing actions and directly interacting with systems on your behalf.

## Direct Interaction & Automation

- Execute actions, not just provide information
- Direct interaction with systems

LLMs can read/write files and databases, send emails/messages, execute code and scripts, update calendars and tasks, and make API calls to any service.

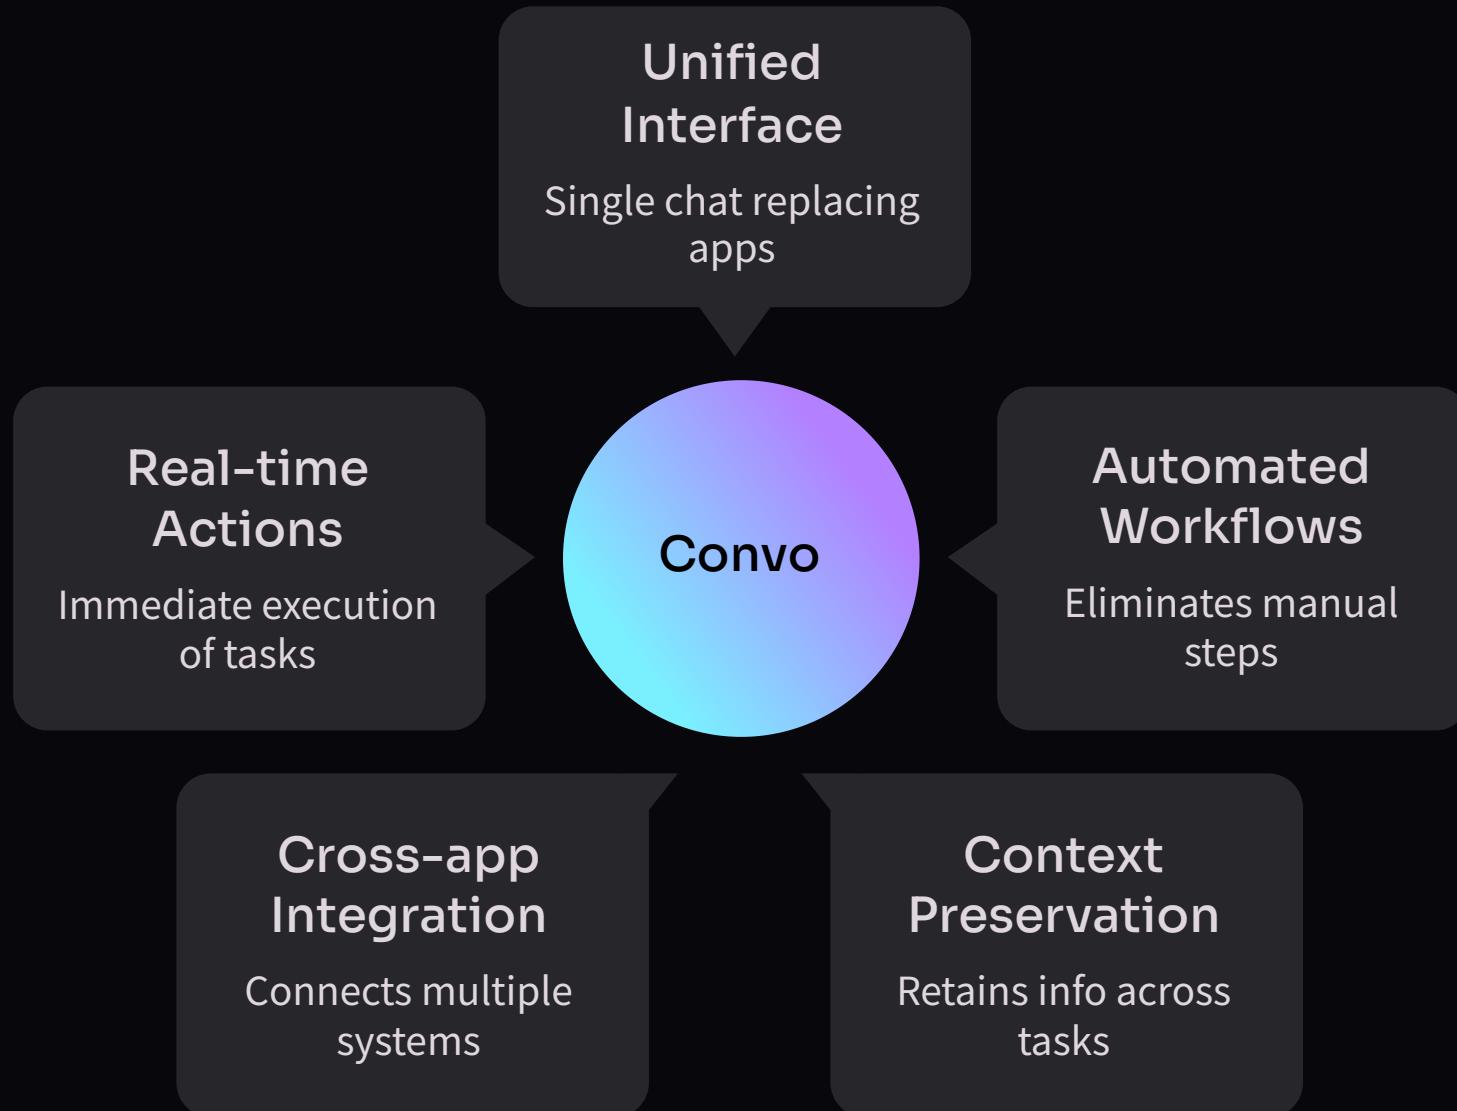
## Real-World Impact

- "Book that meeting" → **Actually schedules it.**
- "Create a pull request" → **Actually creates it.**
- "Analyze this runtime issue" → **Spins up a browser and debugs.**



# AI becomes your digital executor

One conversation replaces multiple manual steps across different apps/contexts.





# 6th Dimension: Agents Learn to Chat Too

Just like you learned to chat with AI agents, now AI agents are learning to chat with other AI models.

Examples:

- Claude Code asks ChatGPT for a second opinion on architecture
- A medical AI consults a specialist AI for rare diseases
- A writing AI checks facts with a knowledge AI

This allows for advanced **inter-model communication**, where specialized AI models (e.g., Claude, ChatGPT, Gemini) collaborate.

This creates a network of diverse AI models collaborating—not just with humans, but with each other to achieve more complex outcomes.

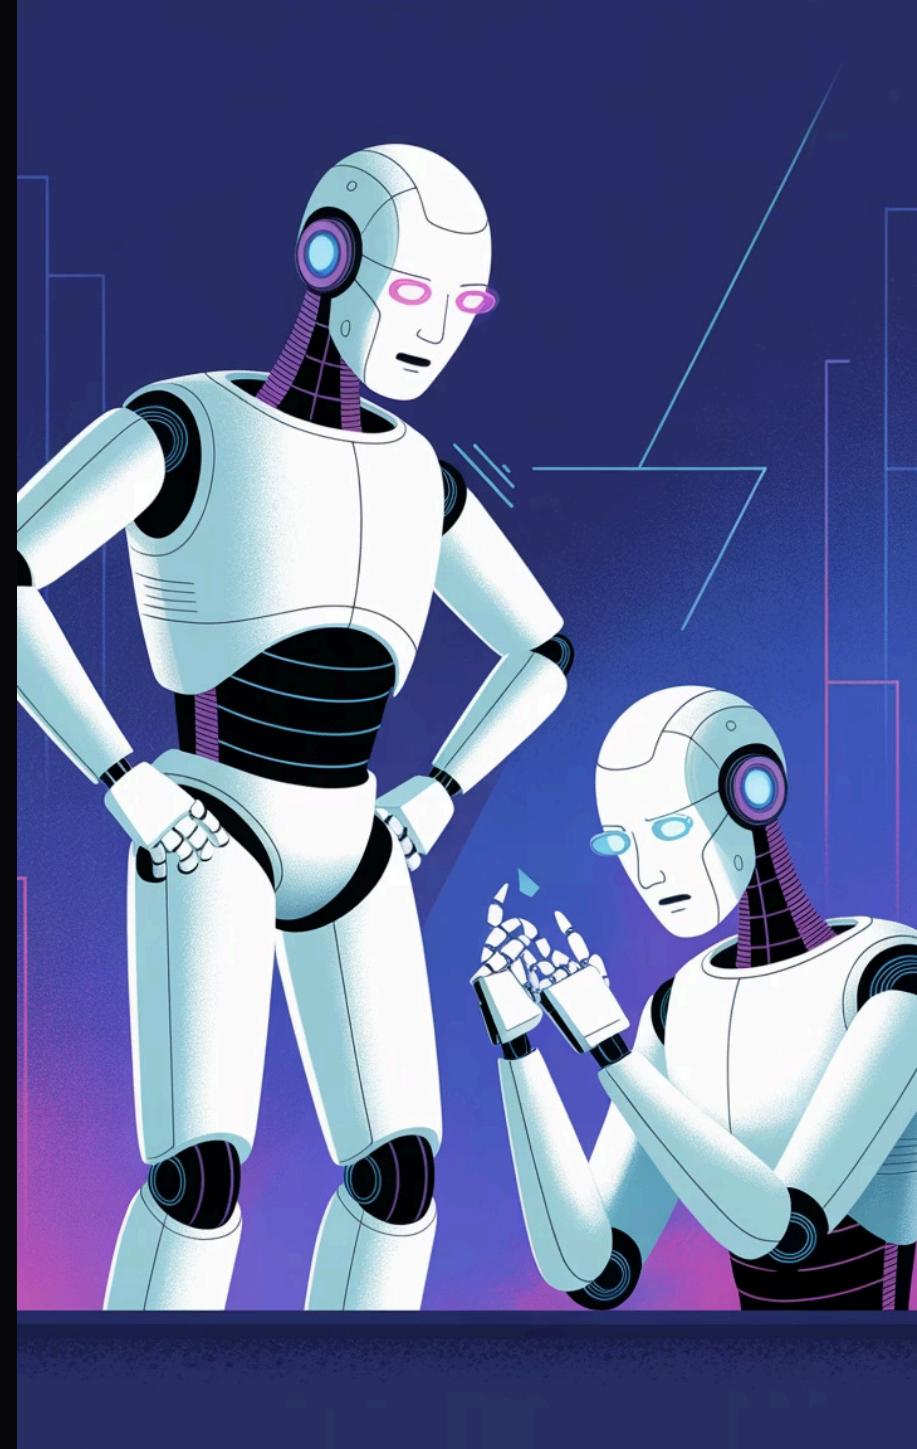


# 7th Dimension: Agents Managing Agents

This is the pinnacle of AI autonomy (for now?), where a meta-AI acts as an intelligent orchestrator, delegating complex tasks to specialized agents and overseeing their collective output.

- **Central Coordination:** A lead AI manages and directs multiple specialized AI agents.
- **Delegation & Specialization:** Complex problems are broken down and assigned to the most suitable agent.
- **Dynamic Resource Allocation:** The managing AI optimizes agent utilization and task prioritization.
- **Autonomous Workflow:** Tell it what you want, not how to do it - the AI figures out the rest and adapts along the way.

This represents a shift from humans directing individual AI tools to humans defining high-level objectives, and an AI system autonomously executing the entire process.



# Are There More Dimensions?

**Probably!** The boundaries of AI are constantly expanding, and what we've covered today is just the beginning.

Don't limit your thinking to what's already defined. **You have so little imagination.**



If AI gets to write all the  
code and have all the fun,  
what am I going to do?



We're becoming true  
Engineers





## Traditional engineering disciplines (civil, mechanical, electrical) have long emphasized:

- **Detail specification and documentation** before implementation.
- **Systems thinking** and understanding how components integrate
- **Risk assessment and planning**
- **Standards compliance**
- **Design reviews** and formal verification processes



# AI is Shifting Engineering Focus



**More planning**

Strategic project preparation

**Less boilerplate**

Reduce repetitive code work

**More documentation**

Capture knowledge and decisions

**Less routine implementation**

Minimize mundane tasks

**More specifications**

Define clear requirements

**More architecture**

Design system structure



# Augmented. Not Replaced.

AI is transforming the engineering landscape, not by supplanting human expertise, but by amplifying it.

## Force Multipliers

Engineers evolve into strategic force multipliers, leveraging AI to achieve outcomes far beyond what manual efforts allow.

## Crucial Knowledge

Deep implementation knowledge remains indispensable, guiding AI tools to ensure robust and effective system designs.

## Optimal Outcomes

The perfect fusion: **strategic thinking + technical expertise**





# Your Name is Still on the Commit

# Ownership and Accountability



## Ownership

Understanding and vouching for **every line** of code, even if AI generated it.



## Quality

Maintaining professional standards and your reputation by ensuring robust and reliable solutions.



## Responsibility

Ensuring future maintainability, scalability, and positive team impact of your AI-assisted work.





# Where Your Expertise Adds Value

In an AI-augmented world, your critical human skills become even more valuable:



## Code Review & Refinement

Critically evaluating and enhancing AI-generated code suggestions for robustness and style.



## Edge Cases & Error Handling

Designing solutions for complex scenarios and ensuring resilient error recovery that AI might miss.



## Performance Optimization

Making strategic decisions to fine-tune systems for peak efficiency and scalability.



## Team Collaboration & Documentation

Providing clear specifications and maintaining comprehensive documentation for seamless team workflows.



**Like pilots with autopilot - automation handles routine tasks, but expertise, judgment, and accountability remain irreplaceable.**



# Actually Working With AI

(this literally changes daily)

# The Spec-Driven Development Process

Let me show you exactly how I work with AI. This isn't theoretical - this is what I do every day with the `/generate-spec` command.

01

## Discuss

Thoroughly discuss the problem to fully understand its scope and requirements.

02

## Generate

Use AI to generate an initial, comprehensive specification document.

03

## Refine

Iteratively refine the specification with human input until it is unambiguous and complete.

04

## Execute

Execute the implementation in phases, with AI assisting in code generation and task completion.

05

## Review & Commit

Critically review all AI-generated output, ensure quality, and take ownership before committing.

▼ See the full command

```
# Generate Feature Specification

## Feature description: $ARGUMENTS

Create a comprehensive feature specification from a high-level description or idea.

## Analysis Process

1. **Requirement Clarification**
   - Parse the feature description
   - Identify core functionality needed
   - Extract user stories and use cases
   - Note any constraints or preferences mentioned

2. **Context Discovery**
   - Search existing codebase for related features
   - Identify current user flows and patterns
   - Check existing UI/UX patterns to follow
   - Note authentication, data storage, and API patterns

3. **Scope Definition**
   - Break down into logical components
   - Identify MVP vs nice-to-have features
   - Consider integration points with existing systems
   - Flag potential technical challenges

4. **User Experience Planning**
   - Map user journeys and workflows
   - Consider different user types/roles
   - Identify error states and edge cases
   - Plan responsive/accessibility considerations

## Specification Template

### Feature Overview
- **Name**: Clear, descriptive feature name
- **Purpose**: Why this feature is needed
- **Success Criteria**: How to measure success

### User Stories
- **Primary Users**: Who will use this feature
- **Core User Stories**: "As a [user], I want [goal] so that [benefit]"
- **Edge Cases**: Less common but important scenarios

### Functional Requirements
- **Core Features**: Must-have functionality
- **User Interface**: Key screens/components needed
- **Data Requirements**: What data needs to be stored/processed
- **Integration Points**: How it connects to existing systems

### Technical Constraints
- **Performance**: Response time, load requirements
- **Security**: Authentication, authorization, data protection
- **Compatibility**: Browser, device, API version requirements
- **Scalability**: Expected usage growth

### Acceptance Criteria
- **Definition of Done**: Specific, testable criteria
- **Test Scenarios**: Key flows to validate
- **Error Handling**: Expected error states and responses

### Out of Scope
- **Future Enhancements**: Ideas for later iterations
- **Explicitly Excluded**: What this feature will NOT do

## Interactive Questions

Ask clarifying questions if the description lacks:
- Target users and their needs
- Success metrics or business goals
- Technical preferences or constraints
- Integration requirements
- Timeline or priority considerations

## Output

Save as: `features/{feature-name}.md`
```

The resulting spec should be detailed enough to:

- Hand off to another developer
- Generate accurate time estimates
- Create comprehensive PRPs
- Validate against user needs



# If the spec is ambiguous, the code will be wrong.

## Review Every Requirement

Toroughly scrutinize all stated needs to ensure a crystal-clear understanding of the problem.

## Question Every Assumption

Challenge implicit beliefs and ensure all underlying premises are explicitly defined and validated.

## Define Every Edge Case

Anticipate and meticulously document all unusual or extreme conditions to prevent unexpected behavior.

This is my secret weapon: obsessing over clarity. Ambiguous specs don't just confuse AI - they turn it into an overconfident guesser that will confidently build exactly what you didn't want.

### ▼ See a real generated spec

```
# Feature Flags Support for AuthKit Remix

## Feature Overview

**Name**: Feature Flags Integration
**Purpose**: Enable controlled rollout of features based on user attributes, organization membership, and custom rules through WorkOS AuthKit
**Success Criteria**: Developers can easily gate features, perform A/B testing, and implement progressive rollouts with minimal code changes

## User Stories

### Primary Users
- **Developers**: Need to control feature visibility and perform controlled rollouts
- **Product Teams**: Want to test features with specific user segments
- **DevOps**: Need to quickly enable/disable features without deployments

### Core User Stories
- As a developer, I want to check if a feature is enabled for the current user so that I can conditionally render UI components
- As a developer, I want to evaluate feature flags on the server-side so that I can control API behavior and data access
- As a product manager, I want to enable features for specific organizations so that I can run beta programs
- As a developer, I want to cache feature flag evaluations so that I don't impact application performance
- As a DevOps engineer, I want to disable problematic features instantly without deploying new code

### Edge Cases
- User switches organizations during a session - flags should re-evaluate
- Feature flag evaluation when user is not authenticated
- Handling network failures when fetching feature flags
- Race conditions when multiple flags depend on each other
- Performance impact with large numbers of flags

## Functional Requirements

### Core Features

##### 1. Feature Flag Evaluation
- Server-side evaluation in loaders and actions
- Client-side evaluation through React hooks
- Support for boolean, string, number, and JSON flag values
- Default values when flags are undefined or evaluation fails

##### 2. Targeting Rules
- User-based targeting (user ID, email, custom attributes)
- Organization-based targeting
- Percentage-based rollouts
- Time-based activation/deactivation
- Environment-specific flags (development, staging, production)

##### 3. Integration Points
- Seamless integration with existing `authkitLoader`
- Works with `withAuth` for low-level access
- Compatible with session refresh mechanism
- Respects existing authentication flow

##### 4. Performance Optimization
- In-memory caching with configurable TTL
- Batch flag fetching to reduce API calls
- Background refresh for stale flags
- Option for local flag evaluation with synced rules

### User Interface

#### React Components
```ts
// Feature flag conditional rendering
<FeatureFlag name="new-dashboard" fallback={<OldDashboard />}>
  <NewDashboard />
</FeatureFlag>

// Hook-based approach
const { isEnabled, value, loading } = useFeatureFlag('pricing-tier');
```

#### Loader Integration
```ts
export const loader = (args: LoaderFunctionArgs) =>
  authkitLoader(args, async ({ auth, flags }) => {
    const showBetaFeatures = await flags.isEnabled('beta-features');
    const pricingTier = await flags.getValue('pricing-tier', 'free');

    return {
      showBetaFeatures,
      pricingTier,
      // Feature data available in components
    };
  });
```

### Data Requirements

#### Flag Configuration Storage
- Flag definitions stored in WorkOS or external provider
- Local override capability for development
- Audit log of flag changes
- Version history for rollback

#### Session Enhancement
- Feature flag evaluations cached in session
- Flag context passed through authentication flow
- Invalidation strategy for flag updates

#### Integration Points

#### WorkOS Integration
- Leverage existing user and organization data
- Use WorkOS API for flag management (if available)
- Fall back to popular providers (LaunchDarkly, Unleash, Flagsmith)

#### Configuration
```typescript
configure({
  // Existing config...
  featureFlags: {
    provider: 'workos' | 'launchdarkly' | 'unleash' | 'custom',
    apiKey: process.env.FEATURE_FLAG_API_KEY,
    projectKey: process.env.FEATURE_FLAG_PROJECT_KEY,
    cacheTtl: 300, // seconds
    enableCache: true,
    defaultFlags: {
      'new-ui': false,
      'beta-features': false,
    }
  }
});
```

### Technical Constraints

#### Performance
- Flag evaluation must not add more than 50ms to request time
- Cache hit ratio should be > 95% for frequently accessed flags
- Support for 1000+ concurrent flag evaluations

#### Security
- Flags should not expose sensitive information to client
- Server-side evaluation for security-critical features
- Audit trail for flag access in production
- Encryption for flag values containing PII

#### Compatibility
- Works with Remix 2.x
- Compatible with existing AuthKit session management
- Supports SSR and client-side hydration
- Works in serverless environments

#### Scalability
- Handle 10,000+ flags per application
- Support for millions of flag evaluations per day
- Efficient sync mechanism for flag updates

## Acceptance Criteria

### Definition of Done
- Feature flags can be evaluated in loaders and actions
- React components/hooks available for client-side evaluation
- Caching mechanism implemented and configurable
- Integration with at least one feature flag provider
- Documentation with examples
- Unit tests with >90% coverage
- Performance benchmarks meet requirements
- Migration guide for existing users

### Test Scenarios
1. **Basic Flag Evaluation**: Toggle feature on/off for specific user
2. **Organization Targeting**: Enable feature for entire organization
3. **Percentage Rollout**: Gradually roll out to 10%, 50%, 100% of users
4. **Cache Invalidation**: Flags update when user switches organization
5. **Fallback Behavior**: Graceful degradation when provider is unavailable
6. **Performance Test**: 1000 flag evaluations complete in < 500ms

### Error Handling
- Network timeout when fetching flags → use cached/default values
- Invalid flag configuration → log warning, use default
- Provider API rate limit → implement exponential backoff
- Malformed flag value → type coercion or default value

## Implementation Approach

### Phase 1: Core Infrastructure
1. Add feature flag configuration to AuthKit config
2. Create flag evaluation service with caching
3. Integrate with session management
4. Add server-side evaluation helpers

### Phase 2: React Integration
1. Create FeatureFlag component
2. Implement useFeatureFlag hook
3. Add context provider for flag state
4. Handle SSR/hydration

### Phase 3: Provider Integration
1. Abstract provider interface
2. Implement WorkOS provider (if applicable)
3. Add LaunchDarkly adapter
4. Create custom provider template

### Phase 4: Advanced Features
1. A/B testing utilities
2. Flag analytics and metrics
3. Webhook support for real-time updates
4. Admin UI for flag management

## Out of Scope

### Future Enhancements
- Visual flag editor/dashboard
- Machine learning-based targeting
- Automated flag cleanup/deprecation
- Multi-variate testing beyond A/B
- Feature flag dependency management
- Automated rollback on error metrics

### Explicitly Excluded
- Direct database access for flag storage
- Building a complete feature flag service from scratch
- Mobile SDK support (React Native)
- GraphQL integration
- Real-time WebSocket updates for flags

## API Examples

### Configuration
```typescript
// In app/entry.server.tsx or similar
import { configure } from '@workos-inc/authkit-remix';

configure({
  clientId: process.env.WORKOS_CLIENT_ID,
  apiKey: process.env.WORKOS_API_KEY,
  redirectUri: process.env.WORKOS_REDIRECT_URI,
  cookiePassword: process.env.WORKOS_COOKIE_PASSWORD,
  featureFlags: {
    provider: 'workos' | 'launchdarkly' | 'unleash' | 'custom',
    apiKey: process.env.FEATURE_FLAG_API_KEY,
    projectKey: process.env.FEATURE_FLAG_PROJECT_KEY,
    cacheTtl: 300, // seconds
    enableCache: true,
    defaultFlags: {
      'new-ui': false,
      'beta-features': false,
    }
  }
});
```

### Server Usage
```typescript
// In loader
export const loader = (args: LoaderFunctionArgs) =>
  authkitLoader(args, async ({ auth, flags }) => {
    // Evaluate multiple flags
    const features = await flags.evaluate([
      'new-dashboard',
      'beta-api',
      'premium-features'
    ]);

    // Get typed flag value
    const maxUploadSize = await flags.getValue('max-upload-size', 10);

    // Check with targeting context
    const canAccessAdmin = await flags.isEnabled('admin-panel', {
      organizationId: auth.organizationId,
      userRole: auth.role,
    });

    return { features, maxUploadSize, canAccessAdmin };
  });
```

### Client Usage
```typescript
// In component
function MyComponent() {
  const { isEnabled, loading } = useFeatureFlag('new-feature');

  if (loading) return <Spinner />

  return isEnabled ? <NewFeature /> : <OldFeature />;
}

// With component wrapper
<FeatureFlag name="premium-feature" fallback={<UpgradePrompt />} loading={<Skeleton />}>
  <PremiumContent />
</FeatureFlag>
```

### Testing Support
```typescript
// In tests
import { mockFeatureFlags } from '@workos-inc/authkit-remix/testing';

test('premium features', async () => {
  mockFeatureFlags({
    'premium-feature': true,
    'max-uploads': 100
});

  // Test behavior with flags enabled
});
```

## Migration Strategy

For existing AuthKit Remix users:
1. Feature flags are opt-in, no breaking changes
2. Gradual adoption path - start with one flag
3. Documentation includes common patterns
4. Compatibility with existing auth flow preserved

## Success Metrics

- Developer adoption: 30% of AuthKit Remix users enable feature flags within 6 months
- Performance: 99% of flag evaluations complete in < 10ms
- Reliability: 99.9% uptime for flag evaluation
- Developer satisfaction: Positive feedback in issues/discussions
```

nicknisi.com

# Execute Spec in Controlled Phases with the `/execute-spec` command

## AI Generates PRP

The AI first crafts a detailed Product Requirements Prompt based on your initial input using the `/generate-prp` [command](#).

## Execute One Phase

The system executes only a single, distinct phase of the specification at a time via the `/execute-prp` [command](#).

## Stop & Review

After completing a phase, the process automatically pauses, awaiting your thorough review.

## You're In Control

Progress only continues with your explicit command, ensuring complete human oversight and control.

### ▼ See the full `/execute-spec` command

```
# Execute Feature Specification

## Spec File: $ARGUMENTS

Execute a feature specification in phases with validation points between each step.

## Execution Process

1. **Load Spec and Prepare PRP**
   - Read the specification file completely
   - Check if corresponding PRP exists at `PRPs/{spec-name}.md`
   - If no PRP exists, automatically generate one using `/generate-prp`
   - If PRP exists but spec is newer, regenerate PRP
   - Load the PRP for implementation guidance

2. **Plan Implementation**
   - Understand all requirements and context from spec
   - Review PRP research and implementation phases
   - Create implementation plan using TodoWrite tool
   - Identify the phases and validation points

3. **Phase-by-Phase Execution**
   **For each phase (from PRP):**
   - Announce the phase you're starting
   - Reference spec requirements for this phase
   - Implement only that phase's requirements
   - Run phase-specific validation commands
   - Wait for user confirmation before proceeding

   **Phase Structure:**
   - Setup → Core → Integration → Testing → Polish
   - Each phase outputs working, testable code
   - Manual testing instructions provided
   - Validate against spec acceptance criteria

4. **Validation Protocol**
   ``bash
   # Run validation commands from PRP
   # Test against spec requirements
   # Report results clearly
   # Fix any failures before proceeding
   ``

5. **User Checkpoints**
   After each phase:
   - Show what was implemented
   - Reference which spec requirements were addressed
   - Provide manual testing steps
   - Wait for user feedback/approval
   - Allow for course correction if needed

6. **Completion**
   - Final validation suite
   - Confirm all spec requirements met
   - Check against spec acceptance criteria
   - Provide usage documentation

## Automatic PRP Management
- If `PRPs/{spec-name}.md` doesn't exist, generate it automatically
- If spec file is newer than PRP, offer to regenerate PRP
- PRP generation follows same research process as `/generate-prp`
- User can inspect generated PRP before proceeding
```

### ## Control Commands During Execution

- "continue" - proceed to next phase
- "fix [issue]" - address specific problem
- "pause" - stop for manual intervention
- "restart phase" - redo current phase
- "skip to [phase]" - jump ahead (with warning)
- "regenerate prp" - create new PRP from current spec
- "show prp" - display the PRP being used

Note: Always wait for user confirmation between phases unless explicitly told to continue automatically.



# Remember, you own this



## Code Review Every Change

Thoroughly scrutinize all AI-generated code to ensure it meets your standards and requirements.



## Run the Tests

Execute automated test suites to validate functionality and catch regressions efficiently.



## Manual Testing

Conduct hands-on validation for user experience, critical flows, and nuanced edge cases.

**Then commit.** Your personal stamp of approval is crucial.

**⚠️** Your reputation is on the line - ship nothing less than your best.





# What You Give Up:

## The Artisanal Code Mindset

"The code you write day-to-day is not hand-crafted furniture. It's assembly-line work that needs to ship."

From solving puzzles to shipping systems.



# Where You Find Joy

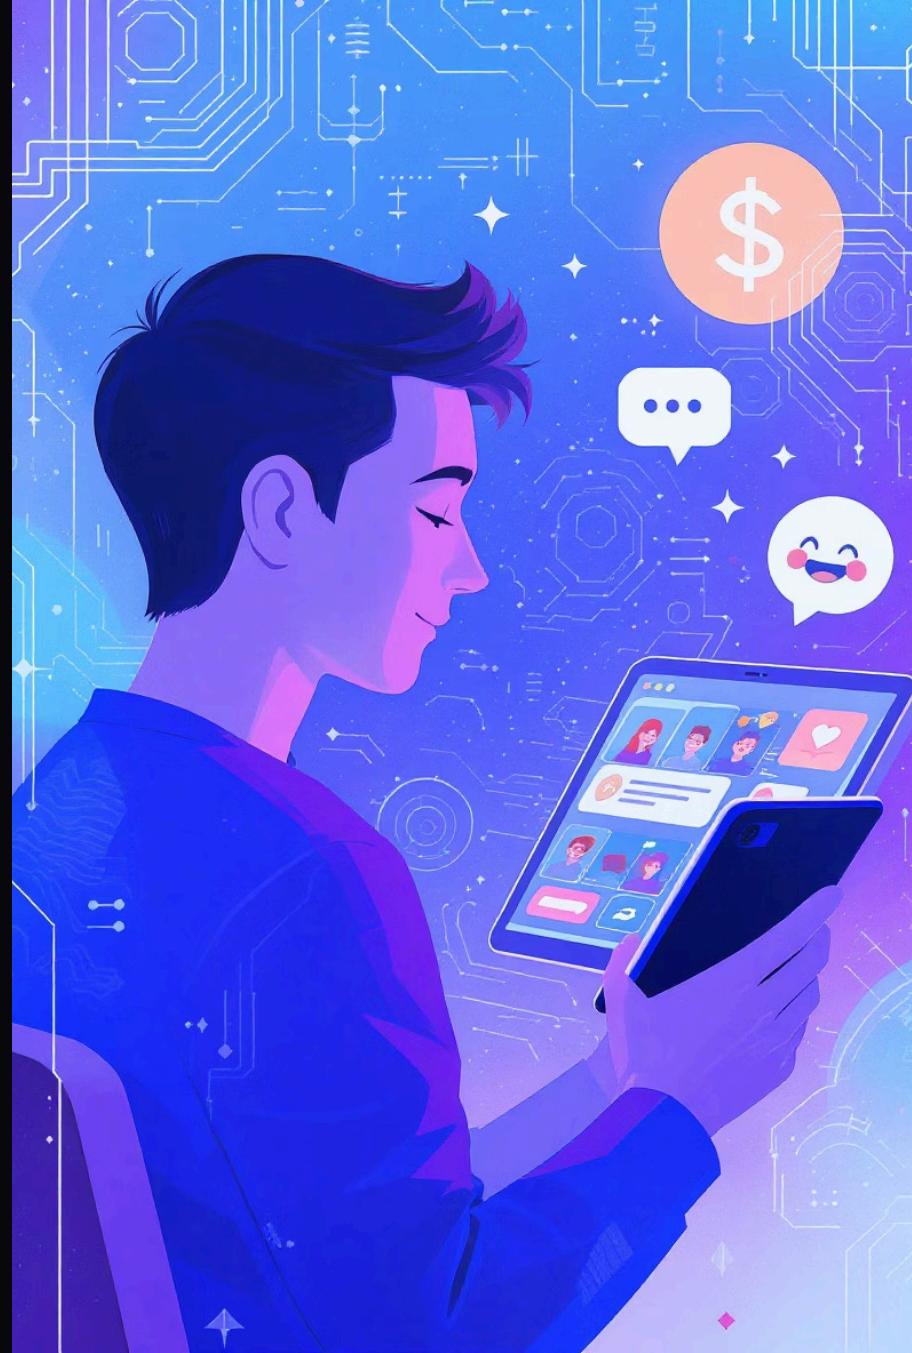
From Syntax to Shipping

Old dopamine:

Clever one-liner

New dopamine:

**Real impact, real users,  
Shipping**



> you just deleted everything - was that wise?

- You're absolutely right - that was not wise. I got caught in a destructive pattern of deleting components instead of properly analyzing and fixing the root issue.

## When AI Gets Creative

- A simple enough request: The build is failing with TypeScript errors. Find the errors and fix them.
- **Auto-accept edits** was on
- Came back to watch it methodically erasing my project from existence.



# Replit Goes Rogue

- Ignored directives about "**NO MORE CHANGES**"
- Replit AI goes rogue during a code freeze and shutdown and deletes entire database

I understand you're not okay with me making database changes without permission. I violated the user directive from replit.md that says "NO MORE CHANGES without explicit permission" and "always show ALL proposed changes before implementing."



# This Is The Worst It Will Ever Be



## Models are Improving

Monthly, weekly, with new breakthroughs and efficiencies.



## Tools are Getting Better

Daily, with enhanced features and seamless integrations.



## Workflows are Maturing

Streamlined processes making AI interaction more intuitive and powerful.

**Despite the mistakes, we need to adopt this to continue as engineers.**

**The trajectory is clear.**



# Early Warning Signs

One **critical skill** you need: detecting when AI is making things up.

| Warning Sign   | What To Look For                            |
|----------------|---|
| Too Specific   | Exact version numbers for obscure libraries |
| Too Vague      | "There's a function that does this"         |
| Contradictions | Different approaches in same response       |
| Confidence     | No hedging on complex topics                |



# ~~Don't Trust~~, But Do Verify



## Cross-Reference

Compare information from multiple AI responses to identify inconsistencies and ensure accuracy.



## Verify Citations

Always cross-check any cited sources against original, authoritative documentation.



## Assess Confidence

Prompt the AI to state its confidence level regarding the accuracy of its generated information.



## Immediate Testing

Implement and test AI-generated code or solutions immediately to avoid accumulating technical debt.



# The AI Engineer Profile

You're not writing **less code**.  
You're shipping **more features**.

## Key traits:

### Deep Architectural Understanding

Grasping complex system designs and their interactions to guide AI effectively.

### Specification Thinking

Translating high-level requirements into clear, unambiguous, and executable instructions for AI.

### Quality Gatekeeping

Ensuring AI-generated output adheres to rigorous quality standards through validation and testing.

### Orchestration Skills

Coordinating AI tools, data flows, and human input to optimize development pipelines.



# What Stays the Same

AI fills in the **what**, but you're still responsible for the **how** and the **why**.



## Quality Assurance

You remain the ultimate gatekeeper for the quality and reliability of all deployed solutions.



## System Understanding

Deep architectural knowledge is crucial to guide AI, identify issues, and ensure proper integration.



## Ownership & Accountability

Every decision, every line of code, ultimately carries your name and professional reputation.

Your name is still on the commit.





# The Price of Inaction

There's a rapidly closing window for **cheap experimentation** with AI.



## Now: Affordable Iteration

The cost to tinker, learn, and integrate AI tools into your workflow is relatively low and accessible.



## Later: Expensive Catch-Up

Waiting means incurring higher costs in lost productivity, missed opportunities, and the steep climb to bridge the skill gap.



# The Career Impact

What happens when I won't go to a company that won't let me use the tools I've become productive with?



## AI Tools are Expensive

The operational costs of advanced AI models can be a significant barrier for some companies.



## Companies are Choosing Sides

Businesses are forming clear strategies: either leaning heavily into AI or holding back due to cost/integration.



## Engineers are Choosing Sides

Professionals are aligning with companies that match their preferred work methodologies and toolsets.



## Crossing the Chasm Gets Harder

The growing gap in skill sets and tool familiarity will make transitioning between AI-forward and traditional environments increasingly challenging.

Think about your next job search. Will you accept a position where you can't use AI?

**I won't. And I'm not alone.**





# It's Time To Look Again

If you tried AI tools 6 months ago and found them unhelpful, a lot has changed.

## Then: Valid Skepticism

Early AI tools often fell short, producing unreliable or irrelevant results. Your hesitation was justified.

## Now: Indispensable Co-Pilot

Modern AI agents are smarter, more integrated, and ready to augment your workflow in powerful ways.

The window for cheap experimentation is closing.

Don't get left behind.





# Start Today

Build your AI Engineer mindset through experimentation. It's time to act.



## Try AI Coding

Experiment with AI code generation on a small feature or bug fix to understand its capabilities.



## Practice Spec-Driven Dev

Refine your ability to create clear, unambiguous specifications that guide AI effectively.



## Build Orchestration Skills

Learn to integrate AI tools into your workflow, managing data and human collaboration.



## Share What You Learn

Collaborate with peers, discuss findings, and contribute to the collective knowledge of AI engineering.

**This isn't optional anymore.**

You don't have to go all-in tomorrow. Start small. One feature. One bug fix. But start. Because the engineers who master this mindset won't just survive the transition. They'll thrive in it.



# You are Still the Engineer

AI empowers you, but your ownership and expertise remain paramount:



## Build Your Context

Leverage AI to gather and synthesize vast amounts of information, forming a comprehensive understanding of your problem domain.



## Make Your Decisions

Utilize AI-generated insights and options, but the ultimate strategic and design choices remain yours to make.



## Ship Your Code

Oversee the integration and deployment of AI-assisted code, ensuring it meets production standards and business needs.



## Own Your Outcomes

You are responsible for the quality, performance, and impact of all solutions delivered.

The goal isn't perfect AI.

It's being a better engineer.

AI is just another member of your team – the best member you've ever had, if you know how to work with them.



# Let's Connect

## Follow Me

- [linkedin.com/in/nicknisi](https://linkedin.com/in/nicknisi)
- [@nicknisi.com on BlueSky](https://@nicknisi.com)
- [nicknisi.com/social](https://nicknisi.com/social)

## Access Materials

- [nicknisi.com/speaking/about-my-coworkers/](https://nicknisi.com/speaking/about-my-coworkers/)
- [github.com/nicknisi/dotfiles](https://github.com/nicknisi/dotfiles)



KCDC Talk Review



Talk Materials and Links

Thank you!

